

LOGISTIC REGRESSION

-APPLIED ANALYTICS-

APM Chapter 12.2, ISLR Chapter 4.3

Lecturer: Darren Homrighausen, PhD

Introductory example

AN INTRODUCTORY EXAMPLE

Suppose we work for a credit card company and we wish to identify people that are likely to default on their credit card debt

We have features (for 10,000 people):

- Student status
- Income
- Balance

Along with their default status:

$$Y = \begin{cases} 1 & \text{if person defaults} \\ 0 & \text{if person doesn't default} \end{cases}$$

Let's look at some plots.

AN INTRODUCTORY EXAMPLE





FIGURE: The red are people without defaults, green are defaults. The '+' are students, the 'Δ' are not students.

AN INTRODUCTORY EXAMPLE



Some comments:

- **Income** doesn't seem to be related to defaults
- **Student status** is also unrelated to defaults, but highly related to **income**
- **Balance** seems to strongly predict default status.

AN INTRODUCTORY EXAMPLE: WHY NOT USE REGRESSION?


Suppose for a moment we only consider balance. Then, we can run a simple linear regression of default status on `balance`

```
Y = rep(0,n)
Y[default == 'Yes'] = 1
out.lm = lm(Y~balance)
summary(out.lm)
```

R will happily do this.

AN INTRODUCTORY EXAMPLE: WHY NOT USE REGRESSION?

Let's plot our data with the fitted values (black) and the training supervisor (red/green)



Not so great..

Logistic regression

LOGISTIC REGRESSION

Logistic regression is a 'generalized linear model' (GLM)

The generalized means that the nature of the supervisor isn't naturally modeled as normal as it is in multiple linear regression:

$$Y = \beta_0 + \sum_{j=1}^p \beta_j x_j + \epsilon$$

(ϵ is modeled as a zero mean, symmetric \rightarrow the residuals should look similar)

If we model **probabilities**, then we are doing **Logistic regression**

LOGISTIC REGRESSION

Suppose for now that $C = 2$, $p_1 = p$, and $p_2 = 1 - p$

(The number of classes, and the probability that Y equals C_1 and C_2 , respectively)

Then logistic regression models the probability as

$$p = \frac{\exp\{\beta_0 + \sum_{j=1}^p \beta_j x_j\}}{1 + \exp\{\beta_0 + \sum_{j=1}^p \beta_j x_j\}}.$$

This gets converted to something that looks like multiple regression via

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{j=1}^p \beta_j x_j$$

(This is the **logit** function)

INTERPRETING THE PARAMETERS

We can interpret the parameters in either linear model:

- **MULTIPLE LINEAR REGRESSION.**

$$f(X) = \beta_0 + \sum_{j=1}^p x_j \beta_j$$

"A one unit change in x_j is associated with a β_j change in the mean of Y , holding all other features constant"

- **LOGISTIC REGRESSION.** $\log\left(\frac{p}{1-p}\right) = \beta_0 + \sum_{j=1}^p \beta_j x_j$

"A one unit change in x_j is associated with a β_j change in the log odds that $Y = C_1$, holding all other features constant"

or

"A one unit change in x_j is associated with a multiplicative $\exp\{\beta_j\}$ change in the odds that $Y = C_1$, holding all other features constant"

In either case, a $\beta_j > 0$ denotes a positive relationship and $\beta_j < 0$ a negative one

ESTIMATION FOR LOGISTIC REGRESSION

With multiple regression, there is a closed form solution:

$$\hat{\beta} = (\mathbb{X}^T \mathbb{X})^{-1} \mathbb{X}^T \mathbb{Y}$$

This was due to estimation via least squares

(This is also the **maximum likelihood estimator** (MLE) under ϵ being normally distributed)

To find the estimator $\hat{\beta}$ for logistic regression, we have to use an iterative maximization technique

Let's begin to look at the output:

ESTIMATING LOGISTIC REGRESSION

```
out.glm = glm(default~balance,family='binomial')
> summary(out.glm)
Call:
glm(formula = default ~ balance, family = "binomial")
Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.2697  -0.1465  -0.0589  -0.0221   3.7589

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.065e+01  3.612e-01  -29.49  <2e-16 ***
balance      5.499e-03  2.204e-04   24.95  <2e-16 ***
...
Number of Fisher Scoring iterations: 8
```

PREDICTIONS: PROBABILITY ESTIMATES

Once we have the $\hat{\beta}$, we can estimate the probabilities via

$$\hat{p} = \frac{\exp\{\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j\}}{1 + \exp\{\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j\}}.$$

EXAMPLE: Suppose we want to predict the probability that someone with a balance of \$1,000 defaults

$$\hat{p}(1000) = \frac{\exp\{-10.65 + 0.0055 * 1000\}}{1 + \exp\{-10.65 + 0.0055 * 1000\}} = 0.0058$$

Maybe look at \$2,000 instead:

$$\hat{p}(2000) = \frac{\exp\{-10.65 + 0.0055 * 2000\}}{1 + \exp\{-10.65 + 0.0055 * 2000\}} = 0.586$$

PREDICTIONS: CLASSIFICATIONS

We need to turn the estimated probabilities:

$$\hat{p} = \frac{\exp\{\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j\}}{1 + \exp\{\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j\}}.$$

into classifications

EXAMPLE: Using the threshold of $1/C = 0.5$,

$$\hat{p}(1000) = \frac{\exp\{-10.65 + 0.0055 * 1000\}}{1 + \exp\{-10.65 + 0.0055 * 1000\}} = 0.0058$$

Thus, a balance of \$1000 would be classified as **no default**

Maybe look at \$2000 instead:


$$\hat{p}(2000) = \frac{\exp\{-10.65 + 0.0055 * 2000\}}{1 + \exp\{-10.65 + 0.0055 * 2000\}} = 0.586.$$

A balance of \$2000 would be classified as **default**

A COMPARISON


Let's plot our data with

- Simple Linear Regression (black)
- Logistic regression (blue)



A COMPARISON

Results of using a cut-off of 0.5



Judging Model Performance

USING CARET

```
trControl      = trainControl(method = 'cv', number = 10)
outLogistic    = train(x = Xtrain, y = Ytrain,
                       method = 'glm', trControl = trControl)
```

Let's get a training/test split as well, resulting in **Xtrain**,
Xtest, **Ytrain**, and **Ytest**

(Use `createDataPartition`)

GETTING PREDICTIONS

```
YhatTestProb = predict(outLogistic, Xtest, type = 'prob')  
> head(YhatTestProb)
```

	Yes	No
1	0.0005292166	0.9994708
2	0.0025568958	0.9974431
3	0.0023380867	0.9976619
4	0.0205685913	0.9794314
5	0.0001117464	0.9998883
6	0.0007992407	0.9992008

```
YhatTest = predict(outLogistic, Xtest, type = 'raw')  
> head(YhatTest)
```

```
[1] No No No No No No
```

```
Levels: Yes No
```

WELL CALIBRATED PROBABILITIES

QUESTION: What is a calibration plot? What are well-calibrated probabilities?


WELL CALIBRATED PROBABILITIES

QUESTION: What is a calibration plot? What are well-calibrated probabilities?

```
calibProbs = calibration(Ytest~YhatTestProb$Yes, cuts = 5)

which(YhatTestProb$Yes > .8)
[1] 987 1004 1216 1735 2059
Ytest[which(YhatTestProb$Yes > .8)]
[1] Yes No No Yes Yes
```

WELL CALIBRATED PROBABILITIES



```
Ytest[which(YhatTestProb$Yes > .8)]  
[1] Yes No No Yes Yes
```

These probability estimates are overall well calibrated

Rare 'Yes' events make the upper probability range unreliable

CONFUSION MATRICES

Using caret, we can report a lot of output

```
> confusionMatrix(data = YhatTest,  
                  reference = Ytest)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Yes	No
Yes	24	12
No	59	2404

Accuracy : 0.9716

...

Kappa : 0.3911

...

Sensitivity : 0.289157

Specificity : 0.995033

...

'Positive' Class : Yes

ADJUSTING THE THRESHOLD

The caret package only allows us to use the $1/C$ threshold:

```
YhatTest = predict(outLogistic, Xtest, type = 'raw')
```

We can generate classifications out of the \hat{p} with a different threshold:

```
YhatTestThresh = ifelse(YhatTestProb$Yes > .2,  
                        'Yes', 'No') %>%  
  as.factor %>%  
  relevel(ref = 'Yes')
```

The **relevel** command is needed to keep the levels consistent

```
> levels(Ytest)  
[1] "Yes" "No"  
> levels(YhatTestThresh)  
[1] "Yes" "No"
```

QUESTION: Do you expect sensitivity to go up or down?

CONFUSION MATRICES

For the threshold 0.2, we get:

```
> confusionMatrix(data = YhatTestThresh,  
                  reference = Ytest)
```

Confusion Matrix and Statistics

	Reference	
Prediction	Yes	No
Yes	50	60
No	33	2356

Accuracy : 0.9628

...

Kappa : 0.4992

...

Sensitivity : 0.60241

Specificity : 0.97517

...


'Positive' Class : Yes

ROC CURVE


We can scan over all the available thresholds with the ROC curve

```
> levels(Ytest)
[1] "Yes" "No"
```

```
rocCurve = roc(relevel(Ytest,'No'), YhatTestProb$Yes)
```



AUC



The area under the curve (AUC) for this classifier is a one number summary of the ROC curve:

```
> rocCurve$auc
```

```
[1] 0.952
```

This would be considered a very good AUC score