

# ADDITIONAL TOPICS ON THE PENALIZATION

## -APPLIED ANALYTICS-

(No book references)

Lecturer: Darren Homrighausen, PhD

# Preamble:

- Sparse matrices provide an efficient way to store data with a lot of zeros
- We can generalize the lasso and ridge regression into the **refitted lasso**

# Sparse matrices

# SPARSE MATRICES

```
load("../data/hiv.rda")
X = hiv.train$x
> X[5:12,1:10]
```

	p1	p2	p3	p4	p5	p6	p7	p8	p9	p10
[1,]	0	0	0	0	0	0	0	0	0	0
[2,]	0	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	0	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	1	0	0
[6,]	0	0	0	0	0	0	0	0	0	0
[7,]	1	0	0	0	0	0	0	0	0	0
[8,]	0	0	0	0	0	0	0	0	0	0

Many zero entries!

# Sparse Matrices

All numbers in **R** take up the same **space**

(Space in this context means RAM aka memory)

```
> print(object.size(0),units='auto')
```

48 bytes

```
> print(object.size(pi),units='auto')
```

48 bytes

**IDEA:** If we can tell **R** in advance which entries are zero ...

- it doesn't need to save those numbers
- nor multiply/add with them

# SPARSE MATRICES

This can be accomplished in several ways in R

I usually use the **Matrix** package

```
library('Matrix')
```

```
Xspar = Matrix(X,sparse=T)
```

# SPARSE MATRICES

Let's take a look at the space difference

```
> print(object.size(X),units='auto')  
1.1 Mb  
> print(object.size(Xspar),units='auto')  
140.7 Kb
```

Pretty substantial! Only 12.1% as large

# SPARSE MATRICES

Lastly, we can create sparse matrices without having the original matrix  $\mathbb{X}$  ever in memory

This is usually done with three vectors of the same length:

- A vector with row numbers
- A vector with column numbers
- A vector with the entry value

```
i = c(1,2,2)
```

```
j = c(2,2,3)
```

```
val = c(pi,1.01,100)
```

```
sparseMat = sparseMatrix(i = i, j = j, x = val,dims=c(4,4))
```

```
regularMat = as(sparseMat,'dgeMatrix')
```



# SPARSE MATRICES

```
> print(sparseMat)
4 x 4 sparse Matrix of class "dgCMatrix"
```

```
[1,] . 3.141593 . .
[2,] . 1.010000 100 .
[3,] . . . .
[4,] . . . .
```

```
> print(regularMat)
4 x 4 Matrix of class "dgeMatrix"
```

```
      [,1]      [,2] [,3] [,4]
[1,]      0 3.141593      0      0
[2,]      0 1.010000    100      0
[3,]      0 0.000000      0      0
[4,]      0 0.000000      0      0
```

# SPARSE MATRICES

Sparse matrices 'act' like regular (**dense**) matrices

They just only keep track of which entries are non zero and perform the operation on these entries

For our purposes, **glmnet** (and other methods) automatically check to see if  $\mathbb{X}$  is a sparse matrix object

This can be a substantial speed/storage savings for large, sparse matrices

**Warning:** be on the look out for your sparse matrix becoming non-sparse!

(Note that the homework has a sparse matrix. You can play around with using the sparse matrix data structure if you like, but for simplicity I didn't include it in the assignment)

# Refitted lasso

# REFITTED LASSO

Since elastic net (for  $\alpha > 0$ ) does both...

- regularization
- model selection

... it tends to produce a solution that

(Here, I'm referring to the elastic net with  $\alpha$  and  $\lambda$  chosen by CV)

- has **too much bias**
- selects too many features


A common approach is to do the following two steps:

1. choose the  $\lambda$  via the 'one-standard-error rule'
2. refit the (unregularized) least squares solution on the selected features

# SOME COMMENTS ABOUT GLMNET AND CV

The function `cv.glmnet` comes with a `plotting` function

```
lassoOut = cv.glmnet(x=X,y=Y,alpha=1)
plot(lassoOut)
```



- The left-most dotted, vertical line occurs at the CV minimum
- The right-most dotted, vertical line is the
  - ▶ largest value of  $\lambda$  ...
  - ▶ such that  $CV(\lambda)$  is within one standard-error of  $CV(\hat{\lambda})$(Note that the vertical bars are at  $2*SE$ )

This is the so called **one-standard-error** rule:  $\hat{\lambda}_{1se}$

# THE ACTIVE SET

Think of the  $p$  features,  $x_j$ , as being indexed from  $1, 2, \dots, p$


Suppose we estimate  $\hat{\beta}_j = 0$ , this is like saying  $x_j$  isn't related to the supervisor

(Hence,  $x_j$  is inactive or not selected)

So, let's collect all the  $j$  such that  $\hat{\beta}_j \neq 0$ , this is often notated  $\hat{S}$

This is known as the **active** or **selected** set

# THE ACTIVE SET



Different values of  $\lambda$  could result in a different active set:  $\hat{\mathcal{S}}(\lambda)$

# REFITTED LASSO

Define the refitted lasso:

1. Get the active set via the '1 standard error rule':  $\hat{S}(\hat{\lambda}_{1se})$
2. Refit with multiple linear regression using only the features in the active set:

$$\text{minimize over } \sum_{i=1}^n \left( Y_i - \left( \beta_0 + \sum_{j \text{ in } \hat{S}(\hat{\lambda}_{1se})} \beta_j X_{ij} \right) \right)^2$$

Report this least squares solution as  $\hat{\beta}_{refitted}$



# REFITTED LOGISTIC LASSO

Note that we can apply this same concept to other methods

In particular, the output from the general elastic net regression  
(assuming  $\alpha > 0$ , of course)

Or to logistic elastic net

(Instead of using least squares, we would use logistic regression)

# Computing the refitted lasso

# REFITTED LASSO

We can train lasso (over a grid of  $\lambda$ ) with **glmnet**

```
require(glmnet)
lassoOut      = cv.glmnet(X ,Y, alpha=1)
```

Now, we need to get  $\hat{S}(\hat{\lambda}_{1se})$

```
betaHatTemp   = coef(lassoOut,s='lambda.1se')[-1]
Srefitted     = which(abs(betaHatTemp) > 1e-16)
```

Lastly, we can get the refitted lasso:

```
Xdf           = as.data.frame(X[,Srefitted])
refittedOut   = lm(Y ~ ., data = Xdf)
betaHatRefitted = coef(refittedOut)
```

If we have a test data set, we can get predictions via

```
XtestDF       = as.data.frame(Xtest[,Srefitted])
YhatTestRefitted = predict(refittedOut, XtestDF)
```

# Postamble:

- Sparse matrices provide an efficient way to store data with a lot of zeros  
(Coercion to sparse matrices saves time and space. Be wary of destroying sparsity, though)
- We can generalize the lasso and ridge regression into the **refitted lasso**  
(The refitted lasso gets the selected features and computes then computes the least squares solution on these selected features)